

# Structured Reactive Programming

Francisco Sant’Anna      Noemi Rodriguez      Roberto Ierusalimsky  
Departamento de Informática — PUC-Rio, Brasil  
{fsantanna,noemi,roberto}@inf.puc-rio.br

## ABSTRACT

.....  
.....  
.....  
.....  
.....  
.....  
.....  
.....  
.....  
.....

## Categories and Subject Descriptors

D.3.3 [Programming Languages]: Language Constructs and Features

## General Terms

Design, Languages

## Keywords

Concurrency, Determinism, Esterel, Imperative, Structured Programming, Synchronous, Reactivity

## 1. INTRODUCTION

Reactive applications interact continuously and in real time with external stimuli from the environment. They represent a wide range of software areas and platforms: from games in powerful desktops, “apps” in capable smart phones, to the emerging internet of things in constrained embedded systems.

Research on special-purpose reactive languages dates back to the early 80’s with the co-development of two complementary styles [3, 17]: The imperative style of Esterel [5] organizes programs with structured control flow primitives, such as sequences, repetitions, and parallelism. The dataflow style of Lustre [12] represents programs as graphs of values, in which a change to a node updates its dependencies automatically.

In recent years, Functional Reactive Programming [20] modernized the dataflow style and became mainstream, deriving a number of languages and libraries, such as Flapjax [15], Rx (from Microsoft), React (from Facebook), and Elm [6]. In contrast, the imperative style of Esterel did not follow this trend and is now confined to the domain of real-time embedded control systems.

As a matter of fact, imperative reactivity is now often associated to the *observer pattern*, typical in object oriented systems, because it heavily relies on side effects [14, 18]. However, short-lived callbacks (i.e., the observers) eliminate any vestige of structured programming, such as support for long-lasting loops and automatic variables [2], which are elementary capabilities of imperative languages. In this sense, callbacks actually disrupt imperative reactivity, becoming “our generation’s goto” [8, 7, 10].

In this work, we revive the imperative style of Esterel, which we now refer as *Structured Reactive Programming (SRP)*. SRP extends the classical hierarchical control constructs of *Structured Programming (SP)* (i.e., concatenation, selection, and repetition [9]) to support continuous interaction with the environment. In practical terms, SRP provides two extensions to SP: an “await <event>” statement that suspends a line of execution until the referred event occurs, keeping all context alive; and parallel constructs that compose multiple lines of execution and make them concurrent. The *await* statement captures the imperative and reactive nature of SRP, restoring sequential execution and support for automatic variables. Parallel compositions<sup>1</sup> allow for multiple *await* statements to coexist, which is necessary to handle concurrent events, common in reactive applications.

We present SRP through the language CÉU [19], a contemporary outlook of imperative reactivity that aims to expand it from the rigid embedded domain. CÉU is based on Esterel and relies on a similar synchronous and deterministic execution model that simplifies the reasoning about concurrency aspects. Our main contribution is a new abstraction mechanism, the *organisms*, that provide a parallel execution body with an object-like interface. In brief, organisms are to SRP like procedures are to SP, i.e., one can abstract a portion of code with a name and manipulate (call) that name from multiple places. There are, however, additional challenges

<sup>1</sup>We use the term *parallel composition* as a synonym for *lexical side-by-side composition*, which does not imply many-core parallel execution.

```

// DECLARATIONS
input <type> <id>;           // external event
event <type> <id>;           // internal event
var <type> <id>;             // variable

// EVENT HANDLING
await <id>;                  // awaits event
emit <id> => <exp>;          // emits event

// COMPOUND STATEMENTS
<...> ; <...> ;              // sequence
if <...> then <...>          // conditional
    else <...> end
loop do <...> end            // repetition
    break                  // (escape repetition)

// PARALLEL COMPOSITIONS
par/and do <...>             // rejoins on both sides
    with <...> end
par/or do <...>              // rejoins on any side
    with <...> end
par do <...>                 // never rejoins
    with <...> end

// C INTEGRATION
_f();                        // C call (prefix '_')
finalize <...>              // finalization
    with <...> end

// ORGANISMS
class <T> with
    <interface>
do
    <body>
end

```

Figure 1: Syntax of Céu.

that apply to organisms:

- Organisms are themselves alive, acting as subprograms with continuous data and control state.
- Organisms are part of a concurrent program that can manipulate and affect their data and execution state.
- Organisms can be dynamically allocated, requiring a memory management model that must also apply to embedded systems.

The rest of the paper is organized as follows: Section 2 presents SRP through Céu, with its underlying synchronous concurrency model, powerful parallel compositions, and the organisms abstraction. Section 3 discusses related work. Section 4 concludes the paper and makes final remarks.

## 2. STRUCTURED REACTIVE PROGRAMMING WITH CÉU

Céu is a concurrent language, in which its lines of execution, known as *trails*, react continuously to external stimuli, all together in synchronous steps. Figure 1 shows a compact reference of Céu.

The introductory example of Figure 2 starts two trails with the **par** construct: the first (lines 4-8) increments variable *v* on every second and prints its value on screen; the second (lines 10-13) resets *v* on every external request to **RESET**. Programs in Céu can access *C* libraries of the underlying platform by prefixing symbols with an underscore (e.g., `_printf(<...>)`).

### 2.1 Synchronous concurrency

```

1 input void RESET; // declares an external event
2 var int v = 0;
3 par do
4     loop do // 1st trail
5         await 1s;
6         v = v + 1;
7         _printf("v = %d\n", v);
8     end
9 with
10    loop do // 2nd trail
11        await RESET;
12        v = 0;
13    end
14 end

```

Figure 2: Introductory example in Céu.

In the synchronous execution model of Céu, a program must react completely to an occurring event before handling the next. A reaction represents a logical instant in which all trails awaiting the occurring event awake and execute, one after the other, until they await again or terminate. During a reaction, the environment is invariant and does not interrupt the running trails<sup>2</sup>. If multiple trails react to the same event, the scheduler employs lexical order to preserve determinism, i.e., the trail that appears first in the source code executes first. To avoid infinite execution for reactions, Céu ensures that all loops have **await** statements on their bodies [19].

As a result of synchronous execution, all operations to variable *v* in Figure 2 are atomic, because reactions to **1s** and **RESET** can never interleave. In contrast, in an asynchronous model with nondeterministic scheduling, the occurrence of **RESET** could preempt the first trail during an increment to *v* (line 6) and reset it (line 12) before printing it (line 7), characterizing a race condition on the variable. The example illustrates the (arguably simpler) reasoning about concurrency aspects under the synchronous execution model.

The synchronous model also empowers SRP with an *orthogonal abortion* construct that simplifies the composition of activities<sup>3</sup>. Orthogonal abortion is the ability to abort an activity from outside it, without affecting the overall consistency of the system (e.g., properly releasing global resources). Figure 3 shows the **par/or** construct of Céu to compose trails and rejoin when either of them terminates, properly aborting the other. The **par/or** is regarded as orthogonal because the composed trails do not know when and how they are aborted (i.e., abortion is external to them). In the figure, each trail has a set of local variables and an execution body that lasts for an arbitrary time. After the **par/or** rejoins, a new set of local variables goes alive, reusing the space from the trails' locals going out of scope.

Abortion in asynchronous languages is challenging [4], because the activity to be aborted might be on a inconsistent state (e.g., holding pending messages or locks). The language runtime needs to wait for the activity to be consistent before aborting it, but what results in two unsatisfactory semantics for a hypothetical **par/or**: either wait to rejoin, making the program unresponsive to incoming events for

<sup>2</sup>The actual implementation enqueues incoming input events to process them in further reactions.

<sup>3</sup>We use the term activity to generically refer to a language's unit of execution (e.g., *thread*, *actor*, *trail*, etc.).

```

par/or do
  // trail A
  <local-variables-A>
  <body-A>
with
  // trail B
  <local-variables-B>
  <body-B>
end
<local-variables-rejoin>
<body-rejoin>

```

**Figure 3:** A `par/or` composes two concurrent trail and rejoins when one terminates, aborting the other.

an arbitrary time; or rejoin immediately and wait in the background, which may cause race conditions with the subsequent code. Immediate rejoin also leads to heap-based allocation (which is discouraged in the context of embedded systems), because activities may coexist with the subsequent code for some time.

As matter of fact, asynchronous languages do not provide effective abortion: CSP only supports a composition operator that “*terminates when all of the combined processes terminate*” [13]; Java’s `Thread.stop` primitive has been officially deprecated [16]; and pthread’s `pthread_cancel` does not guarantee immediate cancellation [1]. Instead, asynchronous activities typically agree a on common protocol to abort each other (e.g., through shared state variables or message passing), increasing coupling among them with concerns that are not directly related to the problem specification.

Synchronous languages, however, provide accurate control over the life cycle of concurrent activities, because in between every reaction, the whole system is idle and consistent [4]. In Section 2.2 we show how CÉU deals with abortion of trails that use stateful resources from the environment, such as file handling and network transmissions.

## 2.2 Parallel compositions

In terms of control structures, SRP extends SP with parallel compositions, allowing applications to handle multiple events concurrently. CÉU provides three parallel constructs, which vary on how they rejoin: a `par/and` rejoins when all trails in parallel terminate; a `par/or` rejoins when any trail in parallel terminates; a `par` never rejoins (even if all trails in parallel terminate).

The example of Figure 4 compares the `par/and` and `par/or` compositions side by side. The code `<...>` represents a complex behavior with any degree of nested compositions. In the `par/and` variation, the behavior repeats every second at minimum, because both sides must terminate before re-executing the loop. In the `par/or` variation, if the behavior does not terminate within 1 second, it is restarted. These SRP archetypes represent, respectively, the *sampling* and *timeout* patterns, which are typical of reactive applications.

The code in Figure 5 relies on hierarchical `par/or` and `par/and` compositions to describe the state machine of a protocol for sensor networks ported to CÉU [11, 19]. The input events `START`, `STOP`, and `RETRANSMIT` (line 1) represent the external interface of the protocol with a client application. The protocol enters the top-level loop and awaits the starting event

<pre> // sampling pattern loop do   par/and do     &lt;...&gt;   with     await 1s;   end end </pre>	<pre> // timeout pattern loop do   par/or do     &lt;...&gt;   with     await 1s;   end end </pre>
--	--

**Figure 4:** The *sampling* and *timeout* patterns using parallel compositions.

```

1  input void START, STOP, RETRANSMIT;
2  loop do
3    await START;
4    par/or do
5      await STOP;
6      with
7        loop do
8          par/or do
9            await RETRANSMIT;
10           with
11             await <rand> s;
12             par/and do
13               await 1min;
14               with
15                 <send-beacon-packet>
16               end
17             end
18           end
19         with
20           <...> // the rest of the protocol
21         end
22       end

```

**Figure 5:** Parallel compositions can describe complex state machines.

(line 3). Once the client application makes a start request, the protocol starts three other trails: one to await the stopping event (line 5); one to periodically transmit a status packet (lines 7-18); and one with the remaining functionality of the protocol (collapsed in line 20). As compositions can be nested, the periodic transmission is another loop that starts two other trails (lines 8-17): one to handle an immediate retransmission request (line 9); and one to await a small random amount of time and transmit the status packet (lines 11-16). The transmission (collapsed in line 15) is enclosed with a `par/and` that takes at least one minute before looping to avoid flooding the network. At any time, the client may request a retransmission (line 9), which terminates the `par/or` (line 8), aborts the ongoing transmission (line 15, if not idle), and restarts the loop (line 7). Also, the client may request to stop the whole protocol (line 5), which terminates the outermost `par/or` and aborts the transmission and all composed trails. In this case, the top-level loop restarts and waits for the next request to start the protocol (line 3), remaining unresponsive to other requests as specified.

The example shows how parallel compositions can describe complex state machines in a structured way, eliminating the use of global state variables for this purpose [19].

### 2.2.1 Finalization

The CÉU compiler tracks the interaction of `par/or` compositions with automatic variables and stateful native *C* functions (e.g., device drivers) to ensure proper trail abortion.

<pre>var _message_t buffer; &lt;fill-buffer-info&gt; _send_enqueue(&amp;buffer); await SEND_ACK;</pre>	<pre>var _message_t buffer; &lt;fill-buffer-info&gt; <b>finalize</b>   _send_enqueue(&amp;buffer) <b>with</b>   _send_dequeue(&amp;buffer); <b>end</b> await SEND_ACK;</pre>
--	--

**Figure 6: Finalization clauses safely release low-level resources.**

Consider the code on the left of Figure 6, which expands the sending trail of Figure 5 (line 15). The `buffer` packet is a local variable whose pointer is passed to native function `_send_enqueue`. The call enqueues the pointer in the radio driver, which holds it up to the occurrence of `SEND_ACK` that acknowledges the packet transmission. In the meantime, the sending trail might be aborted from `STOP` or `RETRANSMIT` requests (lines 5 and 9 of Figure 5), making the packet buffer to go out of scope, and leading to a *dangling pointer* in the radio driver. CéU refuses to compile programs like this and requires *finalization* clauses to accompany unsafe native calls [19]. The code on the right of Figure 6 properly dequeues the packet if the block of `buffer` goes out of scope.

Finalization clauses are fundamental to preserve the orthogonality of `par/or` compositions in SRP.

## 2.3 Organisms

CéU abstracts data and control state into the single concept of organisms. Organisms provide an execution body with multiple lines of execution (control state) as well as object-like interface (data state). A class of organisms describes an interface and a single execution body. The interface exposes public variables, methods, and also internal events. The body can contain any valid code in CéU (including parallel compositions) and starts on instantiation, executing in parallel with the program. Organism instantiation can be either static or dynamic.

We propose a memory model for organisms that eliminates known issues in allocation: *memory leaks*, *dangling pointers*, and the need for *garbage collection*. The model is similar to stack-living local variables of procedures, providing lexical scope and automatic bookkeeping of organisms. We also restrict explicit references to organisms to avoid indirect manipulation.

The example of Figure 7 introduces static organisms through three code chunks:

- The leftmost code (*CODE-1*) blinks two LEDs with different frequencies in parallel and terminates after 1 minute.
- The code in the middle (*CODE-2*) abstracts the blinking LEDs in an organism class and uses two instances to reproduce the same behavior.
- The rightmost code (*CODE-3*) is the equivalent expansion without organisms, which should resemble the original leftmost code.

In *CODE-2*, the `Blink` class (lines 1-9) exposes the `port` and `dt` fields, corresponding to the LED port and blinking period,

respectively. The application then creates two instances, specifying those fields in the constructors (lines 12-15 and 17-20). The constructors also start the organisms bodies (lines 5-8) to run in the background, in parallel with the application. For instance, both instances are already running at the `await 1min` (line 22).

*CODE-3* is semantically equivalent to *CODE-2*, but with the organisms constructors and bodies expanded (lines 11-16 and 20-25). The generated `par/or` makes the instances and the rest of the application (i.e., `await 1min`, in line 28) concurrent. Note the `await FOREVER` statements (lines 17 and 26) that avoid the organisms bodies to terminate the `par/or`. The `_Blink` type (lines 1-4) corresponds to a simple datatype without an execution body.

The main distinction from objects is how organisms can react independently to the environment, i.e., once instantiated they become alive and reactive (hence the name *organisms*). For instance, organisms need not be included in observer lists for events, or rely on the main program to feed their methods with input from the environment. Even though the organisms run independent from the main program, they are still subject to the disciplined synchronous model, which keeps the whole system deterministic.

Another positive aspect of organisms regards to lexical scope and automatic memory management. *CODE-2* uses a `do-end` block (lines 11-23) that limits the scope of the organisms for 1 minute (line 22). During that period, the organisms are accessible (through `b1` and `b2`) and reactive to the environment. Note how the equivalent expansion of *CODE-3* relies on a `par/or` (lines 9-29) to properly abort the organisms bodies after that period (line 28), just before they go out of scope (line 30). In addition, the `par/or` termination properly triggers all active finalization clauses inside the organisms.

Organisms can be explicitly manipulated through their interface variables, methods, and internal events. Here, we illustrate events given that variables and methods work in the same way as in object oriented programming. Figure 8 defines a class `Unit` (lines 1-16) with the event `move` (line 2) that exposes requests to move the unit position. The main program (lines 18-21) creates two units and requests them to move to different positions in a interval of 1 second. The body of the class initializes the current unit's position `pos` and destination position `dst` (lines 4-5). Then, the body enters in a continuous loop (lines 6-15) to handle `move` requests (line 8) while performing the actual moving operation (lines 10-13) in parallel. The `par/or` restarts the loop on every `move` request, which updates the `dst` position. The moving operation can be as complex as needed, for example, using another loop to apply physics over time. The `await FOREVER` (line 13) halts the trail after the move completes. An advantage of event handling over method calls is that they can be composed in the organism body to affect other ongoing operations. In the example, the `await move` aborts and restarts the moving operation, just like the timeout pattern of Figure 4.

### 2.3.1 Dynamic organisms

TODO

<pre> 1 par/or do 2   loop do 3     await 600ms; 4     _toggle(11); 5   end 6 with 7   loop do 8     await 1s; 9     _toggle(12); 10  end 11 with 12   await 1min; 13 end 14 15 /* CODE-1: original blinking */ </pre>	<pre> 1 class Blink with 2   var int port; 3   var int dt; 4 do 5   loop do 6     await (dt)ms; 7     _toggle(port); 8   end 9 end 10 11 do 12   var Blink b1 with 13     this.port = 11; 14     this.dt = 600; 15   end; 16 17   var Blink b2 with 18     this.port = 12; 19     this.dt = 1000; 20   end; 21 22   await 1min; 23 end 24 25 /* CODE-2: blinking organisms */ </pre>	<pre> 1 struct _Blink with 2   var int port; 3   var int dt; 4 end; 5 6 do 7   var _Blink b1, b2; 8 9   par/or do 10    // body of b1 11    b1.port = 0; 12    b1.dt = 2; 13    loop do 14      await (b1.dt)ms; 15      _toggle(b1.port); 16    end 17    await FOREVER; 18  with 19    // body of b2 20    b2.port = 1; 21    b2.dt = 4; 22    loop do 23      await (b2.dt)ms; 24      _toggle(b2.port); 25    end 26    await FOREVER; 27  end 28  with 29    await 1min; 30  end 31 end 32 33 /* CODE-3: organisms expansion */ </pre>
--	--	---

Figure 7: Two blinking LEDs using organisms.

```

1 class Unit with
2   event int move;
3 do
4   var int pos = 0;
5   var int dst = 0;
6   loop do
7     par/or do
8       dst = await this.move;
9     with
10      if dst != pos then
11        <code-to-move-pos-to-dst>
12      end
13      await FOREVER;
14    end
15  end
16 end
17
18 var Unit u1, u2;
19 emit u1.move => 100;
20 await 1s;
21 emit u2.move => 200;

```

Figure 8: Organism manipulation through interface events.

### 2.3.2 References

TODO

## 3. RELATED WORK

TODO

## 4. CONCLUSION

TODO

## 5. REFERENCES

- [1] UNIX man page for pthread\_cancel. man pthread\_cancel.

- [2] A. Adya et al. Cooperative task management without manual stack management. In *ATEC'02*, pages 289–302. USENIX Association, 2002.
- [3] A. Benveniste et al. The synchronous languages twelve years later. In *Proceedings of the IEEE*, volume 91, pages 64–83, Jan 2003.
- [4] G. Berry. Preemption in concurrent systems. In *FSTTCS*, volume 761 of *Lecture Notes in Computer Science*, pages 72–93. Springer, 1993.
- [5] F. Boussinot and R. de Simone. The Esterel language. *Proceedings of the IEEE*, 79(9):1293–1304, Sep 1991.
- [6] E. Czaplicki and S. Chong. Asynchronous functional reactive programming for guis. In *PLDI'13*, pages 411–422, 2013.
- [7] M. de Icaza. Callbacks as our generations' go to statement. <http://tirania.org/blog/archive/2013/Aug-15.html> (accessed in Jul-2014), 2013.
- [8] E. W. Dijkstra. Letters to the editor: go to statement considered harmful. *Communications of the ACM*, 11(3):147–148, 1968.
- [9] E. W. Dijkstra, E. W. Dijkstra, and E. W. Dijkstra. *Notes on structured programming*. Technological University Eindhoven Netherlands, 1970.
- [10] Elm Language Web Site. Escape from callback hell. <http://elm-lang.org/learn/Escape-from-Callback-Hell.elm> (accessed in Jul-2014).
- [11] O. Gnawali et al. Collection tree protocol. In *Proceedings of SenSys'09*, pages 1–14. ACM, 2009.
- [12] N. Halbwachs et al. The synchronous data-flow programming language LUSTRE. *Proceedings of the IEEE*, 79:1305–1320, September 1991.

- [13] C. A. R. Hoare. Communicating sequential processes. *Communications of the ACM*, 21(8):666–677, 1978.
- [14] I. Maier, T. Rompf, and M. Odersky. Deprecating the observer pattern. Technical report, 2010.
- [15] L. A. Meyerovich, A. Guha, J. Baskin, G. H. Cooper, M. Greenberg, A. Bromfield, and S. Krishnamurthi. Flapjax: a programming language for ajax applications. In *ACM SIGPLAN Notices*, volume 44, pages 1–20. ACM, 2009.
- [16] ORACLE. Java thread primitive deprecation. <http://docs.oracle.com/javase/6/docs/technotes/guides/concurrency/threadPrimitiveDeprecation.html> (accessed in Jul-2014), 2011.
- [17] D. Potop-Butucaru et al. The synchronous hypothesis and synchronous languages. In R. Zurawski, editor, *Embedded Systems Handbook*. 2005.
- [18] G. Salvaneschi, G. Hintz, and M. Mezini. Rescala: Bridging between object-oriented and functional style in reactive applications. In *Proceedings of the of the 13th international conference on Modularity*, pages 25–36. ACM, 2014.
- [19] F. Sant’Anna et al. Safe system-level concurrency on resource-constrained nodes. In *Proceedings of SenSys’13*. ACM, 2013.
- [20] Z. Wan and P. Hudak. Functional reactive programming from first principles. *SIGPLAN Notices*, 35(5):242–252, 2000.