

- rely on the synchronous execution model - compositions
- than abstracting over it - finalization, safety - orthogonal
- preemption - simple reasoning
- OS/drivers are stateful

Structured Reactive Programming

Francisco Sant'Anna

Noemi Rodriguez

Roberto Ierusalimsky

Departamento de Informática — PUC-Rio, Brasil

{fsantanna,noemi,roberto}@inf.puc-rio.br

ABSTRACT

Reactive applications interact in real time and continuously with external stimuli from the surrounding environment. They represent a wide range of software areas and platforms: from games in powerful desktops, "apps" in capable smart phones, to the emerging internet of things in constrained embedded systems.

Research on special-purpose reactive languages dates back to the early 80's with the co-development of two complementary styles [1]: The imperative style of Esterel [?] organizes programs with control flow primitives, such as sequences, repetitions, and also parallelism. The dataflow style of Lustre [?] represents programs as graphs of values, in which a change to a node updates its dependencies automatically.

In recent years, Functional Reactive Programming [?] modernized the dataflow style and became mainstream, deriving a number of languages and libraries, such as Flapjax [?], Rx (from Microsoft), React (from Facebook), and Elm [?]. In contrast, the imperative style did not follow this trend and is now confined to the domain of real-time embedded control systems.

As a matter of fact, imperative reactivity is now often associated to the *observer pattern* of object oriented languages, because it heavily relies on side effects over shared data among objects [?, ?]. However, short-lived callbacks (the observers) eliminate any vestige of structured programming, such as loops and automatic variables [?], which are an elementary capability of imperative languages. In this sense, the observer pattern actually disrupts imperative reactivity, becoming "our generation's goto" [?, ?, ?].

In this work, we present a comprehensive set of imperative abstractions for developing structured reactive applications through the programming language CÉU [5]. CÉU is based on Esterel and relies on a similar synchronous and deterministic execution model that simplifies the reasoning about concurrency aspects. CÉU provides a contemporary outlook

of imperative reactivity that aims to expand the application domain of this family with a new abstraction mechanism. In comparison to standard structured programming, CÉU provides three fundamental extensions:

- An `await <evt>` statement to suspend a line of execution until the referred event occurs.
- Parallel constructs (`par`, `par/or`, and `par/and`) to compose multiple lines of execution.
- An abstraction mechanism, named as *organisms*, to reconcile data and control state in a single concept.

The `await` statement captures the imperative and reactive nature of the language, recovering from the inversion of control inherent to the observer pattern, and thus restoring sequential execution and support for automatic variables. Parallel compositions allows for multiple `await` statements to coexist, which is fundamental to handle events concurrently (which reactive applications often require). An organism abstracts parallel and `await` statements and offer an object-like interface that other parts of the program can manipulate.

Permeating all aspects of the programming style, the synchronous execution model restricts but makes possible an underlying model that makes it all reasonable

CONTRIBUTIONS:

determinism

distribution

parallelism

The major contribution of this work

- dynamic synchronous - getting rid of . memory leaks, dangling pointers, and garbage collection

sintaxe (separar section 1 / 2)

DATA+CONTROL and can be dynamically created.

Given concurrency the synchronous model

await inside a sequence or loop form to have multiple sequence allowing for handling multiple events at the same time

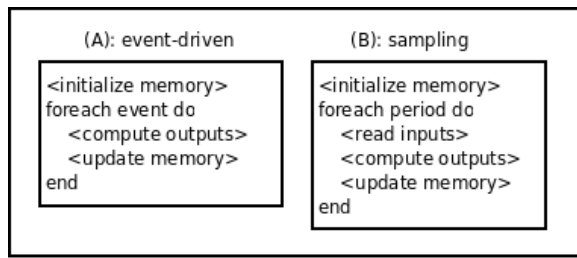


Figure 1: Schedulers for synchronous systems

CÉU has its roots on Esterel and relies on a similar synchronous and deterministic execution model that simplifies the reasoning about concurrency aspects.

these languages shows that imperative can be safe with a reasonable concurrency model and static analysis

SYNCHRONOUS DYNAMIC

sequentiality composibility observer pattern

The rest of the paper is organized as follows: Section 1 reviews the synchronous execution model, which is suitable for reactive programming, and is the base of CÉU. Section 2

1. THE SYNCHRONOUS REACTIVE MODEL

“Reactive systems” are not a new class of problem and have been described by Harel as being “repeatedly prompted by the outside world and their role is to continuously respond to external inputs” citeTODO.harel. In comparison to traditional “transformational systems”, he recognises reactive systems as “particularly problematic when it comes to finding satisfactory methods for behavioral description”. Berry makes a subtle distinction between “interactive” and “reactive” systems [?]:

- Interactive programs interact at their own speed with users or with other programs; from a user point of view a time-sharing system is interactive.
- Reactive programs also maintain a continuous interaction with their environment, but at a speed which is determined by the environment, not by the program itself.

This distinction is fundamental because the different characteristics (i.e., *at the speed of the program* vs *at the speed of the environment*) implies the use of different underlying concurrency models. Overall, *synchronous languages* deal with reactive systems better, while *asynchronous languages*, with interactive systems [?]. Both mentioned authors propose synchronous languages for designing reactive systems (Statecharts [3] and Esterel [2]).

The synchronous execution model is based on the hypothesis that internal computations (*reactions* in this context) run infinitely faster than the rate of events that trigger them. In other words, the input and corresponded output are simultaneous in the synchronous model, because reactions takes no time.

Figure ?? shows two common implementation schemes for synchronous schedulers [1]. In the event-driven scheme, a loop iteration computes outputs for each event occurrence. In the sampling scheme, a loop iteration computes the inputs and outputs for each clock tick. In both cases, each loop iteration represents a logical instant in which the system as a whole reacts synchronously before going to next instant. During a reaction, the environment is invariant, possibly buffering incoming input events to be processed further. Both schemes are compliant with the synchronous hypothesis, in which input and resulting output happen at the same time (in this notion of time as a sequence of discrete events or clock ticks).

The asynchronous execution model is more general and does not make assumptions of implicit synchronization. Each entity in the system (e.g., a thread or actor) is independent from one another and executes at its own speed. In order to coordinate at specific points, they require explicit synchronization primitives (e.g., mutual exclusion or message passing).

For structured reactive programming, we argue that the synchronous model is more appropriate. We discuss two examples with concurrent activities and show that the synchronous model.

to show that We focus on the composition of concurrent activities, given

In our work we use the synchronous model. for concurrency most difficult is composing activities

focus on synchronous vs asynchronous compositions

Figure ?? shows two implementations for an application that blinks two LEDs in parallel with different frequencies. We implemented it in two asynchronous styles and also in CÉU. For *shared memory* concurrency, we used a multithreaded RTOS¹, while for message passing, we used an *occam* variation for Arduino [4].

The intent and syntactic structure of the implementations is very similar: composing in parallel the two blinking activities.

, each blinking a different LED in a loop (simulates a loop with .

but not the semantics

The LEDs should blink together every 3 seconds (*least common denominator* between 600ms and 1s). As we expected, even for such a simple application, the LEDs in the two asynchronous implementations lost synchronism after some time of execution. The CÉU implementation remained synchronized for all tests that we have performed.

The implementations are intentionally naive: they just spawn the activities to blink the LEDs in parallel. The behavior for the asynchronous implementations of the blinking appli-

¹<http://www.chibios.org/dokuwiki/doku.php?id=start>

<pre>// OCCAM-PI PROC main () CHAN SIGNAL s1,s2: PAR tick(600, s1!) toggle(11, s1?) PAR tick(1000, s2!) toggle(12, s2?) :</pre>	<pre>// ChibiOS void thread1 () { while (1) { sleep(600); toggle(11); } } void thread2 () { while (1) { sleep(1000); toggle(12); } } void setup () { create(thread1); create(thread2); }</pre>	<pre>// Ceu par do loop do await 600ms; _toggle(11); end with loop do await 1s; _toggle(12); end end</pre>
---	--	--

Figure 2: Two blinking LEDs in OCCAM-PI, ChibiOS and Céu.

Each line of execution in parallel blinks a LED with a fixed (but different) frequency. (The LEDs are connected to I/O ports 11 and 12.) Every 3 seconds both LEDs should light on together. After a couple of minutes of execution, only the implementation in Céu remains synchronized.

cation is perfectly valid, as the preemptive execution model does not ensure implicit synchronization among activities. In a synchronous language, however, the behavior must be predictable, and loosing synchronism is impossible by design. We used timers in this application, but any kind of high frequency input would also behave nondeterministically in asynchronous systems.

Note that even though the implementations are syntactically similar, with two endless loops in parallel, the underlying execution models between Céu and the two others are antagonistic, hence, the different execution behavior.

Although this application can be implemented correctly with an asynchronous execution model, it circumvents the language style, as timers need to be synchronized in a single thread. Furthermore, it is common to see similar naive blinking examples in reference examples of asynchronous systems², suggesting that LEDs are really supposed to blink synchronized, a guarantee that the language cannot provide (as shown with the examples).

On the one hand, xxx determinism reasoning On the other hand, the synchronous hypothesis does not hold for reactions that have latency (e.g., network communication or algorithm-intensive computations),

On the other hand, execution depends on the order reactions because reaction to events do not interleave in complex

parallelism latency

For reactive systems zzz For highly synchronous systems, the sole synchronization overhead, which is non-existent in

²Example 1 in the RTOS *DuinoOS v0.3*: <http://code.google.com/p/duinos/>.
Example 3 in the occam-based *Concurrency for Arduino v20110201.1855*: <http://concurrency.cc/>.

XXX, may neutralize any gains with parallelism.

Illustrate with two examples: that explore the advantages of synchronous: reasoning in concurrency seamless composition abortion takes *at most* one tick alternatives strong/weak, but *at most* one tick async may take time (unresponsive in the meantime!) even if you simulate with communication, it will take time

CSP has only syntactic support for the *and* semantics p-threads discourages exit Java threads deprecated Exit

∴ determinism and

1-determinism/synchronism 2-composition

communications is directed and takes time (because the receiver may not be waiting)

The synchronous model input => output broadcast possible global consensus possible determinism simpler model

In synchronous systems, communication is instantaneous. The zero-delay property of the synchronous hypothesis guarantees that no time elapses between event announcement and receiving. Also, as communication is via broadcast, all systems parts share the same information all the time. These two characteristics make global data consensus another property of synchronous systems.

hypothesis fail on however, does not apply when the computation involves latency: problem communication takes time either communication or time consuming operations because now, at the speed of the program

Termination and abortion: the real problem with the asynchronous model implies heap-based objects

Two examples: termination / memory / no composition time elapse / non-determinism / analogia do elevador

conclusion: no blind/free/orthogonal composition no deterministic/reproducible execution parallelism / no-forced synch time-consuming / independent

Both seminal works propose synchronous languages state-charts and Esterel so we call the synchrnous reactive model

to contrast with interactive or the new asynchronous reactive model

A recent

important distinction from interactive is only the "environment speed" a huge difference as it implies...

This definition is not the same as today:

The meaning of Reactive

synchronous reactive vs asynchronous reactive

berry classification bash asynchronous compositions PAR/OR
not really parallel

diagrama com as duas formas de implementacao mostrar que
em ambas, apenas um evento eh tratado "of course that re-
actions can trigger time-consuming operation that last mul-
tiple reactions, but it is important to recognize that the
programmer is consious about that..."

Berry classified reactive applications as... From this, he de-
nied threads (ada, now threads) and also CSP-like languages
where communication is point-wise and takes time (csp, now
erlang) asynchronous execution or asynchronous communi-
cation (with latency)

in terms of communication - broadcast (single global vi-
sion of an event) - p2p (loose simultaneity when simulating
broadcast

2. STRUCTURED REACTIVE PROGRAMMING

- await, sequence, vars - parallel, patterns, ortogonal - fi-
nalize for C integration and keep orthogonality - organisms,
static, dynamic. do T;

3. RELATED WORK

- asynchronous langs - Esterel + descendants - CRP - simula
- FRP

interactive vs reactive asynchronous vs synchronous dataflow
vs control dynamic vs static

imperative - sequential (eliminates state machines) - better
resource control - less abstract

dataflow

exemplo data melhor vs control melhor

The synchronous concurrency model...

We show composibility, sequential, imperative safety Then,
we extend synchronous reactive programming with dynamic

control vs data reactivity

Adding dynamic state and references to SRP is a significant
contribution. However the presentation assumes the audi-

ence already appreciates the advantages of SRP over stan-
dard imperative semantics. Most observers will not notice
the difference at all and think this is just a weird syntactic
sugar for standard multithreaded programming. SRP needs
to be better explained and justified, perhaps by showing how
it simplifies the examples.

CÉU is a Esterel-based reactive language that targets con-
strained embedded platforms. Relying on a deterministic se-
mantics, it provides safe shared-memory concurrency among
lines of execution. CÉU introduces a stack-based execution
policy for internal events which enables advanced control
mechanisms considering the context of embedded systems,
such as exception handling and a limited form of coroutines.
The conjunction of shared-memory concurrency with inter-
nal events allows programs to express dependency among
variables reliably, reconciling the control and dataflow reac-
tive styles in a single language.

Categories and Subject Descriptors

D.3.1 [Programming Languages]: Formal Definitions and
Theory; D.3.3 [Programming Languages]: Language Con-
structs and Features

General Terms

Design, Languages

Keywords

Concurrency, Dataflow, Determinism, Embedded Systems,
Esterel, Synchronous, Reactivity

4. REFERENCES

- [1] A. Benveniste et al. The synchronous languages twelve
years later. In *Proceedings of the IEEE*, volume 91,
pages 64–83, Jan 2003.
- [2] F. Boussinot and R. de Simone. The Esterel language.
Proceedings of the IEEE, 79(9):1293–1304, Sep 1991.
- [3] D. Harel. Statecharts: A visual formalism for complex
systems. *Science of Computer Programming*,
8(3):231–274, June 1987.
- [4] C. L. Jacobsen et al. Concurrent event-driven
programming in occam-pi for the Arduino. In *CPA '11*,
volume 68, pages 177–193, June 2011.
- [5] F. Sant'Anna et al. Safe system-level concurrency on
resource-constrained nodes. In *Proceedings of
SenSys'13*. ACM, 2013.